

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/220171387>

# An Efficient Parallel Algorithm for Matrix-Vector Multiplication.

Article in *International Journal of High Speed Computing* · March 1995

DOI: 10.1142/S0129053395000051 · Source: DBLP

CITATIONS

80

READS

1,588

3 authors, including:



**Bruce Hendrickson**

Sandia National Laboratories

154 PUBLICATIONS 10,548 CITATIONS

[SEE PROFILE](#)



**Robert Leland**

Sandia National Laboratories

58 PUBLICATIONS 3,195 CITATIONS


[SEE PROFILE](#)

# SANDIA REPORT

SAND92-2765 • UC-405

Unlimited Release

Printed March 1993



\*8570032\*

SANDIA NATIONAL  
LABORATORIES  
TECHNICAL LIBRARY

## An Efficient Parallel Algorithm for Matrix—Vector Multiplication

CONFIDENTIAL

Bruce Hendrickson, Robert Leland, Steve Plimpton

Prepared by  
Sandia National Laboratories  
Albuquerque, New Mexico 87185 and Livermore, California 94550  
for the United States Department of Energy  
under Contract DE-AC04-76DP00789

Issued by Sandia National Laboratories, operated for the United States Department of Energy by Sandia Corporation.

**NOTICE:** This report was prepared as an account of work sponsored by an agency of the United States Government. Neither the United States Government nor any agency thereof, nor any of their employees, nor any of their contractors, subcontractors, or their employees, makes any warranty, express or implied, or assumes any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represents that its use would not infringe privately owned rights. Reference herein to any specific commercial product, process, or service by trade name, trademark, manufacturer, or otherwise, does not necessarily constitute or imply its endorsement, recommendation, or favoring by the United States Government, any agency thereof or any of their contractors or subcontractors. The views and opinions expressed herein do not necessarily state or reflect those of the United States Government, any agency thereof or any of their contractors.

Printed in the United States of America. This report has been reproduced directly from the best available copy.

Available to DOE and DOE contractors from  
Office of Scientific and Technical Information  
PO Box 62  
Oak Ridge, TN 37831

Prices available from (615) 576-8401, FTS 626-8401

Available to the public from  
National Technical Information Service  
US Department of Commerce  
5285 Port Royal Rd  
Springfield, VA 22161

NTIS price codes  
Printed copy: A03  
Microfiche copy: A01

## An Efficient Parallel Algorithm for Matrix-Vector Multiplication

Bruce Hendrickson, Robert Leland and Steve Plimpton  
Sandia National Laboratories  
Albuquerque, NM 87185

### Abstract

**Abstract.**

The multiplication of a vector by a matrix is the kernel computation of many algorithms in scientific computation. A fast parallel algorithm for this calculation is therefore necessary if we are to make full use of the new generation of parallel supercomputers. This paper presents a high performance, parallel matrix-vector multiplication algorithm that is particularly well suited to hypercube multiprocessors. For an  $n \times n$  matrix on  $p$  processors, the communication cost of this algorithm is  $O(n/\sqrt{p} + \log(p))$ , independent of the matrix sparsity pattern. The performance of the algorithm is demonstrated by employing it as the kernel in the well-known NAS conjugate gradient benchmark, where a run time of 6.09 seconds was observed. This is the best published performance on this benchmark achieved to date using a massively parallel supercomputer.

**Key words.** matrix-vector multiplication, parallel computing, hypercube, conjugate gradient method

**AMS(MOS) subject classification.** 65Y05, 65F10

**Abbreviated title.** Parallel Matrix-Vector Multiplication.

---

This work was supported by the Applied Mathematical Sciences program, U.S. Department of Energy, Office of Energy Research, and was performed at Sandia National Laboratories, operated for the U.S. Department of Energy under contract No. DE-AC04-76DP00789.

**1. Introduction.** The multiplication of a vector by a matrix is the kernel computation in many linear algebra algorithms, including, for example, the popular Krylov methods for solving linear and eigen systems. Recent improvements in such methods, coupled with the increasing use of massively parallel computers, require the development of efficient parallel algorithms for matrix–vector multiplication. This paper describes such an algorithm. Although the method works on all parallel architectures, it is particularly well suited to machines with hypercube interconnection topology, for example the Intel iPSC/860 and the nCUBE 2.

The algorithm described here was developed independently in connection with research on efficient methods of organizing parallel many-body calculations (see [5]). We subsequently learned that our algorithm is very similar in structure to a parallel matrix-vector multiplication algorithm described in [4]. We have, nevertheless, chosen to present our algorithm because it improves upon that in [4] in several ways: First, we specify how to overlap communication and computation and thereby reduce the overall run time. Second, we show how to map the blocks of the matrix to processors in a novel way which improves the performance of a critical communication operation on current hypercube architectures. And third, we consider the actual use of the algorithm within the iterative conjugate gradient solution method and show how in this context a small amount of redundant computation can be used to further reduce the communication requirements. By integrating these improvements we have been able to achieve significantly better performance on a well known benchmark than has been previously possible with a massively parallel machine.

A very attractive property of the new algorithm is that its communication operations are independent of the sparsity pattern of the matrix, making it applicable to all matrices. For an  $n \times n$  matrix on  $p$  processors, the cost of the communication is  $O(n/\sqrt{p} + \log(p))$ . However, many sparse matrices exhibit structure which allows for other algorithms with even lower communication requirements. Typically this structure arises from the physical problem being modeled by the matrix equation and manifests itself as the ability to reorder the rows and columns to obtain a nearly block–diagonal matrix, where the  $p$  diagonal blocks are about equally sized, and the number of matrix elements not in the blocks is small. This structure can also be expressed in terms of the size of the separator of the graph describing the nonzero structure of the matrix. Our algorithm is clearly not optimal for such matrices, but there are many contexts where the matrix structure is not helpful (e.g. dense matrices, random matrices), or the effort required to identify the structure is too large to justify. It is these settings in which our algorithm is most appropriate and provides high performance.

This paper is structured as follows. In the next section we describe the algorithm and its communication primitives. In §3 we present refinements and improvements to the basic algorithm, and develop a performance model. In §4 we apply the algorithm to the NAS conjugate gradient benchmark problem to demonstrate its utility. Conclusions are drawn in §5.

**2. A parallel matrix–vector multiplication algorithm.** Iterative solution methods for linear and eigen systems are one of the mainstays of scientific computation. These methods involve repeated matrix–vector products or *matvecs* of the form  $y_i = Ax_i$  where the new iterate,  $x_{i+1}$ , is generally

some simple function of the product vector  $y_i$ . To sustain the iteration on a parallel computer, it is necessary that  $x_{i+1}$  be distributed among processors in the same fashion as the previous iterate  $x_i$ . Hence, a good matvec routine will return a  $y_i$  with the same distribution as  $x_i$  so that  $x_{i+1}$  can be constructed with a minimum of data movement. Our algorithm respects this distribution requirement.

We will simplify notation and consider the parallel matrix-vector product  $y = Ax$  where  $A$  is an  $n \times n$  matrix and  $x$  and  $y$  are  $n$ -vectors. The number of processors in the parallel machine is denoted by  $p$ , and we assume for ease of exposition that  $n$  is evenly divisible by  $p$  and that  $p$  is an even power of 2. It is fairly straightforward to relax these restrictions.

Let  $A$  be decomposed into square blocks of size  $(n/\sqrt{p}) \times (n/\sqrt{p})$ , each of which is assigned to one of the  $p$  processors, as illustrated by Fig. 1. We introduce the Greek subscripts  $\alpha$  and  $\beta$  running from 0 to  $\sqrt{p} - 1$  to index the row and column ordering of the blocks. The  $(\alpha, \beta)$  block of  $A$  is denoted by  $A_{\alpha\beta}$  and owned by processor  $P_{\alpha\beta}$ . The input vector  $x$  and product vector  $y$  are also conceptually divided into  $\sqrt{p}$  pieces indexed by  $\beta$  and  $\alpha$  respectively. Given this block decomposition, processor  $P_{\alpha\beta}$  must know  $x_\beta$  in order to compute its contribution to  $y_\alpha$ . This contribution is a vector of length  $n/\sqrt{p}$  which we denote by  $z_{\alpha\beta}$ . Thus  $z_{\alpha\beta} = A_{\alpha\beta}x_\beta$ , and  $y_\alpha = \sum_\beta z_{\alpha\beta}$  where the sum is over all the processors sharing row block  $\alpha$  of the matrix.




Fig. 1. Structure of matrix product  $y = Ax$ .

**2.1. Communication primitives.** Our algorithm requires three distinct patterns of communication. The first of these is an efficient method for summing elements of vectors owned by different processors, and is called a *fold* operation in [4]. We will use this operation to combine contributions to  $y$  owned by the processors that hold a block row of  $A$ . The fold operation is sketched in Fig. 2 for communication among processors with the same block row index  $\alpha$ . Each processor begins the fold operation with a vector  $z_{\alpha\beta}$  of length  $n/\sqrt{p}$ . The operation requires  $\log_2(\sqrt{p})$  stages, halving the length of the vectors involved at each stage. Within each stage, a processor first divides its vector  $z$  into two equal sized subvectors,  $z_1$  and  $z_2$ , as denoted by  $(z_1|z_2)$ . One of these subvectors is sent to another

processor, while the other processor sends back its contribution to the subvector which remained. The received subvector is summed element-by-element with the retained subvector to finish the stage. At the conclusion of the fold, each processor has a unique, length  $n/p$  portion of the fully summed vector. We denote this subvector with Greek superscripts, hence  $P_{\alpha\beta}$  owns portion  $y^{\alpha\beta}$ . The fold operation requires no redundant floating point operations, and the total number of values sent and received by each processor is  $n/\sqrt{p} - n/p$ .

```

Processor  $P_{\alpha\beta}$  knows  $z_{\alpha\beta} \in \mathbb{R}^{n/\sqrt{p}}$ 
 $z := z_{\alpha\beta}$ 
For  $i = 0, \dots, \log_2(\sqrt{p}) - 1$ 
     $(z_1 | z_2) = z$ 
     $P_{\alpha\beta'} := P_{\alpha\beta}$  with  $i^{\text{th}}$  bit of  $\beta$  flipped
    If bit  $i$  of  $\beta$  is 1 Then
        Send  $z_1$  to processor  $P_{\alpha\beta'}$ 
        Receive  $w_2$  from processor  $P_{\alpha\beta'}$ 
         $z_\beta := z_2 + w_2$ 
    Else
        Send  $z_2$  to processor  $P_{\alpha\beta'}$ 
        Receive  $w_1$  from processor  $P_{\alpha\beta'}$ 
         $z := z_1 + w_1$ 
 $y^{\alpha\beta} := z$ 
Processor  $P_{\alpha\beta}$  now owns  $y^{\alpha\beta} \in \mathbb{R}^{n/p}$ 

```

**Fig. 2.** The fold operation for processor  $P_{\alpha\beta}$  as part of block row  $\alpha$ .

In the second communication operation each processor knows some information that must be shared among all the processors in a column. We use a simple algorithm called *expand* [4], that essentially uses the inverse communication pattern of the fold operation. The expand operation is outlined in Fig. 3 for communication between processors with the same column index  $\beta$ . Each processor in the column begins with a subvector of length  $n/p$ , and when the operation finishes all processors in the column know all  $n/\sqrt{p}$  values in the union of their subvectors. At each step in the operation a processor sends all the values it knows to another processor and receives that processor's values. These two subvectors are concatenated, as indicated by the “|” notation. As with the fold operation, only a logarithmic number of stages are required, and the total number of values sent and received by each processor is  $n/\sqrt{p} - n/p$ .

The optimal implementation of the fold and expand operations depends on the machine topology and various hardware considerations, *e.g.* the availability of multiport communication. There are, however, efficient implementations on most architectures. On hypercubes, for example, these operations can be implemented using only nearest neighbor communication if the blocks in each row and column of the matrix are owned by a subcube with  $\sqrt{p}$  processors. On meshes, if the blocks of the matrix are

```

Processor  $P_{\alpha\beta}$  knows  $y^{\beta\alpha} \in \mathbb{R}^{n/p}$ 
 $z := y^{\beta\alpha}$ 
For  $i = \log_2(\sqrt{p}) - 1, \dots, 0$ 
     $P_{\alpha',\beta} := P_{\alpha\beta}$  with  $i^{\text{th}}$  bit of  $\alpha$  flipped
    Send  $z$  to processor  $P_{\alpha',\beta}$ 
    Receive  $w$  from processor  $P_{\alpha',\beta}$ 
    If bit  $i$  of  $\alpha$  is 1 Then
         $z := w|z$ 
    Else
         $z := z|w$ 
 $y_\alpha := z$ 
Processor  $P_{\alpha\beta}$  now knows  $y_\alpha \in \mathbb{R}^{n/\sqrt{p}}$ 

```

**Fig. 3.** The expand operation for processor  $P_{\alpha\beta}$  as part of block column  $\beta$ .

mapped in the natural way to a square grid of processors, then the fold and expand operations can be implemented efficiently [9].

The third communication operation in our matvec algorithm requires each processor to send a message to the processor owning the transpose portion of the matrix, *i.e.*  $P_{\alpha\beta}$  sends to  $P_{\beta\alpha}$ . Since we want row and column communication to be efficient for the fold and expand operations, this transpose communication can be difficult to implement efficiently. This is because a large number of messages must travel to architecturally distant processors, so the potential for message congestion is great. We have devised an optimal, congestion-free algorithm for this operation on hypercubes which is discussed in §3.1. Implementations of our matvec algorithm on other architectures may benefit from a similarly tailored transpose algorithm. However, even if congestion is unavoidable, the length of the message in the transpose communication step of our matvec algorithm is about  $\sqrt{p}$  less than the volume of data exchanged in the fold and expand steps. Consequently, the transpose messages can be delayed by  $O(\sqrt{p})$  without changing the overall scaling of the algorithm.

**2.2. The matrix–vector multiplication algorithm.** We can now present our algorithm for computing  $y = Ax$  in Fig. 4. Further details and enhancements are presented in the following section. All the numerical operations in the algorithm are performed in steps (1) and (2). First, in step (1), each processor performs the local matrix–vector multiplication involving the portion of the matrix it owns. These values are summed within processor rows in step (2) using the fold operation from Fig. 2, after which each processor owns  $n/p$  of the values of  $y$ . Unfortunately, the values owned by processor  $P_{\alpha\beta}$  are just a subvector of  $y_\alpha$ , whereas to perform the next matvec  $P_{\alpha\beta}$  must know all of  $y_\beta$ . This is accomplished in steps (3) and (4). In step (3), each processor exchanges its subsegment of  $y$  with the processor owning the transpose block of the matrix. After the transposition, the values of  $y_\beta$  are distributed among the processors in column block  $\beta$  of  $A$ . The expand operation among these processors gives each of them all of  $y_\beta$ , so the result is distributed as required for a subsequent matvec.



We note that at this level of detail, the algorithm is identical to the one described in [4] for dense matrices, but as we discuss in the next section, the details of steps (1), (2) and (3) are different and result in a more efficient overall algorithm.

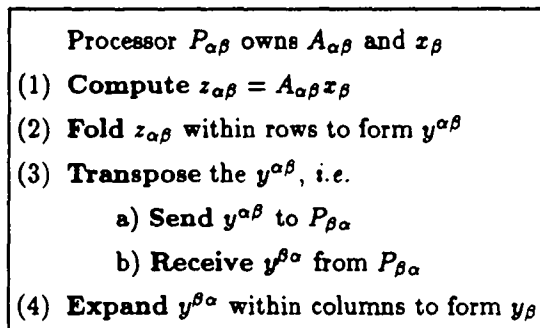


Fig. 4. Parallel matrix-vector multiplication algorithm for processor  $P_{\alpha\beta}$ .

### 3. Algorithmic details and refinements.

**3.1. Transposition on parallel computers.** The expand and fold primitives used in the matvec algorithm are most efficient on a parallel computer if rows and columns of the matrix are mapped to subsets of processors that allow for fast communication. On a hypercube a natural subset is a subcube, while on a 2-D mesh rows, columns or submeshes are possible. Unfortunately, such a mapping can make the transpose operation inefficient since it requires communication between processors that are architecturally distant. Modern parallel computers use *cut-through routing* so that a single message can be transmitted between non-adjacent processors at nearly the same speed as if it were sent between adjacent processors. Nevertheless, if multiple messages are simultaneously trying to use the same wire, all but one of them must be delayed. Hence machines with cut through routing can still suffer from serious message congestion.

On a hypercube the scheme for routing a message is usually to compare the bit addresses of the sending and receiving processors and flip the bits in a fixed order (and transmit along the corresponding channel) until the two addresses agree. On the nCUBE 2 and Intel iPSC/860 hypercubes the order of comparisons is from lowest bit to highest, a procedure known as *dimension order routing*. Thus a message from processor 1001 to processor 0100 will route from 1001 to 1000 to 1100 to 0100. The usual scheme of assigning matrix blocks to processors uses low order bits to encode the column number and the high order bits to encode the row number. Unfortunately, dimension order routing on this mapping induces congestion since messages from all the  $\sqrt{p}$  processors in a row route through the diagonal processor. A similar bottleneck occurs with mesh architectures where the usual routing scheme is to move within a row before moving within a column. Fortunately, the messages being transposed in our algorithm are shorter than those in the fold and expand operations by a factor of  $\sqrt{p}$ . So even if congestion delays the transpose messages by  $\sqrt{p}$ , the overall communication scaling of the algorithm will not be affected.

On a hypercube, a different mapping of matrix blocks to processors can avoid transpose congestion

altogether. With this mapping we still have nearest neighbor communication in the fold and expand operations, but now the transpose operation is as fast as sending and receiving a single message of length  $n/p$ . Consider a  $d$ -dimensional hypercube where the address of each processor is a  $d$ -bit string. For simplicity we assume that  $d$  is even. The row block number  $\alpha$  is a  $d/2$ -bit string, as is the column block number  $\beta$ . For fast fold and expand operations, we require that the processors in each row and column form a subcube. This is assured if any set of  $d/2$  bits in the  $d$ -bit processor address encode the block row number and the other  $d/2$  bits encode the block column number. Now consider a mapping where the bits of the block row and block column indices of the matrix are interleaved in the processor address. For a 64-processor hypercube (with 3-bit row and column addresses for the 8x8 blocks of the matrix) this means the 6-bit processor address would be  $r_2c_2r_1c_1r_0c_0$  where the three bits  $r_2r_1r_0$  encode the block row index and  $c_2c_1c_0$  encodes the block column index.

Note that in this mapping each row of blocks and column of blocks of the matrix still resides on a subcube of the hypercube, so the expand and fold operations can be performed optimally. However, the transpose operation is now contention-free as demonstrated by the following theorem. Although the proof assumes a routing scheme where bits are flipped in order from lowest to highest, a similar contention free mapping is possible for any fixed routing scheme as long as row and column bits are forced to change alternately.

**THEOREM 3.1.** *Consider a hypercube using dimension order routing, and map processors to elements of an array in such a way that the bit-representations of a processor's row number and column number are interleaved in the processor's bit-address id. Then the wires used when each processor sends a message to the processor in the transpose location in the array are disjoint.*

*Proof.* Consider a processor  $P$  with bit-address  $r_b c_b r_{b-1} c_{b-1} \dots r_0 c_0$ , where the row number is encoded with  $r_b \dots r_0$ , and the column number with  $c_b \dots c_0$ . The processor  $P^T$  in the transpose array location will have with bit-address  $c_b r_b c_{b-1} r_{b-1} \dots c_0 r_0$ . Under dimension order routing, a message is transmitted in as many stages as there are bits, flipping bits in order from right to left to generate a sequence of intermediate patterns. After each stage, the message will have been routed to the intermediate processor denoted by the current intermediate bit pattern. The wires used in routing the message from  $P$  to  $P^T$  are those that connect two processors whose patterns occur consecutively in the sequence of intermediate patterns. After  $2k$  stages, the intermediate processor will have the pattern  $r_b c_b \dots r_k c_k c_{k-1} r_{k-1} \dots c_0 r_0$ . The bits of this intermediate processor are a simple permutation of the original bits of  $P$  in which the lowest  $k$  pairs of bits have been swapped. Also, after  $2k - 1$  stages, the values in the bit positions  $2k$  and  $2k - 1$  are equal.

Now consider another processor  $P' \neq P$ , and assume that the message being routed from  $P'$  to  $P^T$  uses the same wire employed in step  $i$  of the transmission from  $P$  to  $P^T$ . Denote the two processors connected by this wire by  $P_1$  and  $P_2$ . Since they differ in bit position  $i$ ,  $P_1$  and  $P_2$  can only be encountered consecutively in the transition between stages  $i - 1$  and  $i$  of the routing algorithm. Either  $i - 1$  or  $i$  is even, so a simple permutation of pairs of bits of  $P$  must generate either  $P_1$  or  $P_2$ ; say  $P_*$ . Similarly, the same permutation applied to  $P'$  must also yield either  $P_1$  or  $P_2$ ; say  $P'_*$ . If  $P_* = P'_*$  then  $P = P'$  which is a contradiction. Otherwise, both  $P_1$  and  $P_2$  must appear after an odd number of

stages in one of the routing sequences. If  $i$  is odd then bits  $i$  and  $i + 1$  of  $P$  must be equal, and if  $i$  is even then bits  $i$  and  $i - 1$  of  $P$  are equal. In either case,  $P_1 = P_2$  which again implies the contradiction that  $P = P'$ .  $\square$

**3.2. Overlapping computation and communication.** If a processor is able to both compute and communicate simultaneously, then the algorithm in Fig. 4 has the shortcoming that once a processor has sent a message in the fold or expand operations, it is idle until the message from its neighbor arrives. This can be alleviated in the fold operation in step (2) of the algorithm by interleaving communication with computation from step (1). Rather than computing all the elements of  $z_{\alpha\beta}$  before beginning the fold operation, we should compute just those that are about to be sent. Then whichever values will be sent in the next pass through the fold loop get computed between the send and receive operations in the current pass. In the final pass, the values that the processor will keep are computed. In this way, the total run time is reduced on each pass through the fold loop by the minimum of the message transmission time and the time to compute the next set of elements of  $z_{\alpha\beta}$ .

**3.3. Balancing the computational load.** The discussion above has concentrated on the communication requirements of our algorithm, but an efficient algorithm must also ensure that the computational load is well balanced across the processors. For our algorithm, this requires balancing the computations within each local matvec. If the region of the matrix owned by a processor has  $m'$  nonzeros, the number of floating point operations (flops) required for the local matvec is  $2m' - n/\sqrt{p}$ . These will be balanced if  $m' \approx m/p$  for each processor, where  $m$  is the total number of nonzero elements in the matrix. For dense matrices or random matrices in which  $m \gg n$ , the load is likely to be balanced. However for matrices with some structure it may not be. For these problems, Ogielski and Aiello have shown that randomly permuting the rows and columns gives good balance with high probability [8]. A random permutation has the additional advantage that zero values encountered when summing vectors in the fold operation are likely to be distributed randomly among the processors.

Most matrices used in real applications have nonzero diagonal elements. We have found that when this is the case, it may be advantageous to force an even distribution of these among processors and to randomly map the remaining elements. This can be accomplished by first applying a random symmetric permutation to the matrix. This preserves the diagonal while moving the off-diagonal elements. The diagonal can now be mapped to processors to match the distribution of the  $y^{\alpha\beta}$  subsegment that each processor owns. The contribution of the diagonal elements can then be computed in between the send and receive operations in the transpose communication, saving either the transpose transmission time or the diagonal computation time, whichever is smaller.

**3.4. Complexity model.** The algorithm described above can be implemented to require the minimal  $2m - n$  flops to perform a matrix-vector multiplication, where  $m$  is the number of nonzeros in the matrix. Some of these flops will occur during the calculation of the local matvecs, and the rest during the fold summations. We make no assumptions about the data structure used on each processor to compute its local matrix-vector product. This allows for the implementation of whatever algorithm works best on the particular hardware. If we assume the computational load is balanced by using the

techniques described in §3.3, the time to execute these floating point operations should be very nearly  $(2m - n)T_{\text{flop}}/p$ , where  $T_{\text{flop}}$  is the time required for a single floating point operation.

The algorithm requires  $\log_2(p) + 1$  read/write pairs for each processor, and a total communication volume of  $n(2\sqrt{p} - 1)$  floating point numbers. Accounting for the natural parallelism in the communication operations, the effective communication volume is  $n(2\sqrt{p} - 1)/p$ . Unless the matrix is very sparse, the computational time required to form the local matvec will be sufficient to hide the transmission time in the fold operation, as discussed in §3.2. We will assume that this is the case. Furthermore, we will assume that the transpose transmission time can be hidden with computations involving the matrix diagonal, as described in §3.3. The effective communication volume therefore reduces to  $n(\sqrt{p} - 1)/p$ . The total run time,  $T_{\text{total}}$  can now be expressed as

$$(1) \quad T_{\text{total}} = \frac{2m - n}{p}T_{\text{flop}} + (\log_2(p) + 1)(T_{\text{send}} + T_{\text{receive}}) + \frac{n(\sqrt{p} - 1)}{p}T_{\text{transmit}},$$

where  $T_{\text{flop}}$  is the time to execute a floating point operation,  $T_{\text{send}}$  and  $T_{\text{receive}}$  are the times to initiate a send and receive operation respectively, and  $T_{\text{transmit}}$  is the transmission time per floating point value. This model will be most accurate if message contention is insignificant, as it is with the mapping for hypercubes described in §3.1.

**4. Application to the Conjugate Gradient algorithm.** To examine the efficiency of our parallel matrix-vector multiplication algorithm, we used it as the kernel of a conjugate gradient (CG) solver. A version of the CG algorithm for solving the linear system  $Ax = b$  is depicted in Fig. 5. There are a number of variants of the basic CG method; the one we present is a slightly modified version of the algorithm given in the NAS benchmark [1, 3] discussed later. In addition to the matrix-vector multiplication, the inner loop of the CG algorithm requires three vector updates (of  $x$ ,  $r$  and  $p$ ), as well as two inner products (forming  $\gamma$  and  $\rho'$ ).

An efficient parallel implementation of the CG algorithm should divide the workload evenly among processors while keeping the cost of communication small. Unfortunately, these goals are in conflict because when the vector updates are distributed, the inner product calculations require communication among all the processors. In addition, if the algorithm in Fig. 5 is implemented in parallel, each processor must know the value of  $\alpha$  before it can update  $r$  to compute  $\rho'$  and hence  $\beta$ . The calculation of  $\gamma = p^T y$ , the distribution of  $\gamma$ , and the calculation of  $\rho' = r^T r$  can actually be condensed into two global operations because the first two operations can be accomplished simultaneously with a binary exchange algorithm. However these global operations are still very costly. One way to reduce the communication load of the algorithm is to modify it as shown in Fig. 6.

This modified algorithm is algebraically equivalent to the original, but instead of updating  $r$  and then calculating  $r^T r$ , the new algorithm exploits the identity  $r_{i+1}^T r_{i+1} = (r_i - \alpha y)^T (r_i - \alpha y) = r_i^T r_i - \alpha y^T r_i + \alpha^2 y^T y$ , as suggested by Van Rosendale [10]. The values of  $\gamma$ ,  $\phi$  and  $\psi$  can be summed with a single global communication, essentially halving the communication time required outside the matvec routine. In exchange for this communication reduction, there is a net increase of one inner product calculation since  $\phi = y^T r$  and  $\psi = y^T y$  must now be computed, but  $\rho' = r^T r$  need not

```

 $x := 0$ 
 $r := b$ 
 $p := b$ 
 $\rho := r^T r$ 
For  $i=1, \dots$ 
     $y := Ap$ 
     $\gamma := p^T y$ 
     $\alpha := \rho / \gamma$ 
     $x := x + \alpha p$ 
     $r := r - \alpha y$ 
     $\rho' := r^T r$ 
     $\beta := \rho' / \rho$ 
     $\rho := \rho'$ 
     $p := r + \beta p$ 

```

**Fig. 5.** A conjugate gradient algorithm.

be calculated explicitly. Since the vectors are distributed across all the processors, this requires an additional  $2n/p$  floating point operations by each processor in order to avoid a global communication. Whether this is a net gain depends upon the relative sizes of  $n$  and  $p$ , as well as the cost of flops and communication on a particular machine, but since communication is typically much more expensive per unit than computation, the modified algorithm should generally be faster. For the nCUBE 2, the machine used in this study, we estimate that this recasting of the algorithm is worthwhile when  $n \leq 5 \times 10^5$ .

This restructuring of the CG algorithm can in principle be carried further to hide more of the communication cost of the linear solve. That is, by repeatedly substituting for the residual and search vectors  $r$  and  $p$  we can express the current values of these vectors in terms of their values  $k$  steps previously. (General formulas for this process are given in [7].) By proper choice of  $k$  it is possible to completely hide the global communication in the CG algorithm. Unfortunately this leads to a serious loss of stability in the CG process which is expensive to correct [6]. We therefore recommend only limited application of this restructuring idea.

The vector and scalar operations associated with CG fit conveniently between steps (3) and (4) of the matrix-vector multiplication algorithm outlined in Fig. 4. At the end of step (3) the product vector  $y$  is distributed across all  $p$  processors, and it is trivial to achieve the identical distribution for  $x$ ,  $r$  and  $p$ . Now all the vector updates can proceed perfectly in parallel. At the end of the CG loop, the vector  $p$  can be shared through an expand operation within columns and hence the processors will be ready for the next matvec. The resulting algorithm is sketched in Fig. 7.

We implemented a double precision version of this algorithm in C on the 1024 processor nCUBE 2 hypercube at Sandia's Massively Parallel Computing Research Laboratory. The resulting code was

```

 $x := 0$ 
 $r := b$ 
 $p := b$ 
 $\rho := r^T r$ 
For  $i=1, \dots, n$ 
     $y := Ap$ 
     $\gamma := p^T y$ 
     $\phi := y^T r$ 
     $\psi := y^T y$ 
     $\alpha := \rho / \gamma$ 
     $\rho' := \rho - \alpha \phi + \alpha^2 \psi$ 
     $\beta := \rho' / \rho$ 
     $\rho := \rho'$ 
     $x := x + \alpha p$ 
     $r := r - \alpha y$ 
     $p := r + \beta p$ 

```

**Fig. 6.** A modified conjugate gradient algorithm.

tested on the well-known NAS parallel benchmark problem proposed by researchers at NASA Ames [1, 3]. The benchmark uses a conjugate gradient iteration to approximate the smallest eigenvalue of a random, symmetric matrix of size 14,000, with an average of just over 132 nonzeros in each row. The benchmark requires 15 calls to the conjugate gradient routine, each of which involves 25 passes through the innermost loop containing the matvec.

This benchmark problem has been addressed by a number of different researchers on several different machines [2]. A common theme in this previous work has been the search for some exploitable structure within the benchmark matrix. Since arbitrary restructuring of the matrix is permitted by the benchmark rules as a pre-processing step, the computational effort expended in this search for structure is not counted in the benchmark timings.

In contrast, our algorithm is completely generic and does not require any special structure in the matrix. The communication operations are entirely independent of the zero/nonzero pattern of the matrix, and the only advantage of reordering would be to lessen the load on the most heavily burdened processor. Because the benchmark matrix diagonal is dense, we did partition the diagonal across all processors, as described in §3.3. Otherwise, we accepted the matrix as given, and made no effort to exploit structure.

Our implementation solved the benchmark problem in 6.09 seconds, which compares quite favorably with all other published results on massively parallel machines [3]. For comparison, the recently published times for the 128 processor iPSC/860 and 32K CM-2 are 8.61 and 8.8 seconds respectively, which is more than 40% longer than our result. Although this problem is highly unstructured, our

```

Processor  $P_{\mu\nu}$  owns  $A_{\mu\nu}$ 
 $x, r, p, b, y \in \mathbb{R}^{n/p}, z_\mu, p_\nu \in \mathbb{R}^{n/\sqrt{p}}$ 
 $x := 0$ 
 $r := b$ 
 $p := b$ 
 $\bar{\rho} := r^T r$ 
Sum  $\bar{\rho}$  over all processors to form  $\rho$ 
Expand  $p$  within columns to form  $p_\nu$ 
For  $i = 1, \dots$ 
    Compute  $z_\mu = A_{\mu\nu} p_\nu$ 
    Fold  $z_\mu$  within rows to form  $y^{\mu\nu}$ 
    Transpose  $y^{\mu\nu}$ , i.e.
        Send  $y^{\mu\nu}$  to  $P_{\nu\mu}$ 
        Receive  $y := y^{\nu\mu}$  from  $P_{\nu\mu}$ 
     $\bar{\gamma} := p^T y$ 
     $\bar{\phi} := y^T r$ 
     $\bar{\psi} := y^T y$ 
    Sum  $\bar{\gamma}, \bar{\phi}$  and  $\bar{\psi}$  over all processors to form  $\gamma, \phi$  and  $\psi$ 
     $\alpha := \rho / \gamma$ 
     $\rho' := \rho - \alpha \phi + \alpha^2 \psi$ 
     $\beta := \rho' / \rho$ 
     $\rho := \rho'$ 
     $x := x + \alpha p$ 
     $r := r + \alpha y$ 
     $p := r + \beta p$ 
    Expand  $p$  within columns to form  $p_\nu$ 

```

Fig. 7. A parallel CG algorithm for processor  $P_{\mu\nu}$ .

C code achieves about 250 Mflops, which is about 12% of the raw speed achievable by running pure assembly language BLAS on each processor without any communication.

**5. Conclusions.** We have presented a parallel algorithm for matrix-vector multiplication, and shown how this algorithm can be used very effectively within the conjugate gradient algorithm. The communication cost of this algorithm is independent of the zero/nonzero structure of the matrix and scales as  $n/\sqrt{p}$ . Consequently, the algorithm is most appropriate for matrices in which structure is either difficult or impossible to exploit. This is clearly the case for dense and random matrices, and it is also true more generally for sparse matrices in many contexts. For example, our algorithm could serve as an efficient black-box routine for prototyping sparse matrix linear algebra algorithms or could be embedded in a sparse matrix library where few assumptions about matrix structure can be made.

On the NAS conjugate gradient benchmark, an nCUBE 2 implementation of this algorithm runs more than 40% faster than any other reported algorithm running on any massively parallel machine.

The particular mapping we employ for hypercubes is likely to be of independent interest. This mapping ensures that rows and columns of the matrix are owned entirely by subcubes, and that with cut-through routing the transpose operation can be performed without message contention. This mapping has already proved useful for parallel many-body calculations [5], and is probably applicable to other linear algebra algorithms.

**Acknowledgements.** We are indebted to David Greenberg for assistance in developing the hypercube transposition algorithm in §3.1.

## REFERENCES

- [1] D. H. BAILEY, E. BARSZCZ, J. T. BARTON, D. S. BROWNING, R. L. CARTER, L. DAGUM, R. A. FATOOHI, P. O. FREDERICKSON, T. A. LASINSKI, R. S. SCHREIBER, H. D. SIMON, V. VENKATAKRISHNAN, AND S. K. WEERATUNGA, *The NAS parallel benchmarks*, Intl. J. Supercomputing Applications, 5 (1991), pp. 63–73.
- [2] D. H. BAILEY, E. BARSZCZ, L. DAGUM, AND H. D. SIMON, *NAS parallel benchmark results*, in Proc. Supercomputing '92, IEEE Computer Society Press, 1992, pp. 386–393.
- [3] D. H. BAILEY, J. T. BARTON, T. A. LASINSKI, AND H. D. SIMON, EDITORS, *The NAS parallel benchmarks*, Tech. Rep. RNR-91-02, NASA Ames Research Center, Moffett Field, CA, January 1991.
- [4] G. C. FOX, M. A. JOHNSON, G. A. LYZENGA, S. W. OTTO, J. K. SALMON, AND D. W. WALKER, *Solving problems on concurrent processors: Volume 1*, Prentice Hall, Englewood Cliffs, NJ, 1988.
- [5] B. HENDRICKSON AND S. PLIMPTON, *Parallel many-body calculations without all-to-all communication*, Tech. Rep. SAND 92-2766, Sandia National Laboratories, Albuquerque, NM, December 1992.
- [6] R. W. LELAND, *The Effectiveness of Parallel Iterative Algorithms for Solution of Large Sparse Linear Systems*, PhD thesis, University of Oxford, Oxford, England, October 1989.
- [7] R. W. LELAND AND J. S. ROLLETT, *Evaluation of a parallel conjugate gradient algorithm*, in Numerical methods in fluid dynamics III, K. W. Morton and M. J. Baines, eds., Oxford University Press, 1988, pp. 478–483.
- [8] A. T. OGIELSKI AND W. AIELLO, *Sparse matrix computations on parallel processor arrays*, SIAM J. Sci. Stat. Comput., 14 (1993). To appear.
- [9] R. VAN DE GEIJN, *Efficient global combine operations*, in Proc. 6th Distributed Memory Computing Conf., IEEE Computer Society Press, 1991, pp. 291–294.
- [10] J. VAN ROSENDALE, *Minimizing inner product data dependencies in conjugate gradient iteration*, in 1983 International conference on parallel processing, H. J. Siegel et al., eds., IEEE, 1983, pp. 44–46.



# EXTERNAL DISTRIBUTION:

R. W. Alewine  
DARPA/RMO  
1400 Wilson Blvd.  
Arlington, VA 22209

Firooz Allahadi  
US Air Force Weapons Lab  
Nuclear Technology Branch  
Kirtland, AFB, NM 87117-6008

Michael W. Allen  
The Wall Street Journal  
1233 Regal Row  
Dallas, TX 75247

M. Alme  
Alme and Associates  
6219 Bright Plume  
Columbia, MD 21044

Charles E. Anderson  
Southwest Research Institute  
PO Drawer 28510  
San Antonio, TX 78284

Dan Anderson  
Ford Motor Co., Suite 1100  
Village Place  
22400 Michigan Ave.  
Dearborn, MI 48124

Andy Arenth  
National Security Agency  
Savage Road  
Ft. Meade, MD 20755  
Attn: C6

Greg Astfalk  
Convex Computer Corp.  
3000 Waterview Parkway  
PO Box 833851  
Richardson, TX 75083-3851

Susan R. Atlas  
TMC, MS B258  
Center for Nonlinear Studies  
Los Alamos National Labs  
Los Alamos, NM 87545

L. Auslander  
DARPA/DSO  
1400 Wilson Blvd.  
Arlington, VA 22209

D. M. Austin  
Army High Per Comp. Res. Cntr.  
University of Minnesota  
1100 S. Second St.  
Minneapolis, MN 55415

Scott Baden  
University of California, San Diego  
Dept. of Computer Science  
9500 Gilman Drive  
Engineering 0114  
La Jolla, CA 92091-0014

F. R. Bailey  
MS200-4  
NASA Ames Research Center  
Moffett Field, CA 94035

D. H. Bailey  
NAS Applied Research Branch  
NASA Ames Research Center  
Moffett Field, CA 94035

Raymond A. Bair  
Molecular Science Resch. Cntr.  
Pacific NW Laboratory  
Richland, WA 99352

R. E. Bank  
Dept. of Mathematics  
Univ. of CA at San Diego  
La Jolla, CA 92093

Ken Bannister  
US Army Ballistic Res. Lab  
Attn: SLCBR-IB-M  
Aberdeen Prov. Grnd., MD 21005-5066

Edward Barragy  
Dept. ASE/EM  
University of Texas  
Austin, TX 78712

E. Barszcz  
NAS Applied Research Branch  
NASA Ames Research Center  
Moffett Field, CA 94035

W. Beck  
AT&T Pixel Machines, 4J-214  
Crawfords Corner Rd.  
Homdel, NJ 07733-1988

David J. Benson  
Dept. of AMES R-011  
Univ. of California at San Diego  
La Jolla, CA 92093

Myles R. Berg  
Lockheed, O/62-30, B/150  
1111 Lockheed Way  
Sunnyvale, CA 94089-3504

Stephan Bilyk  
US Army Ballist. Resch. Lab  
SLCBB-TB-AMB  
Aberdeen Proving Ground, MD  
21005-5066

Rob Bisseling  
Shell Research B.V.  
Postbus 3003  
1003 AA Amsterdam  
The Netherlands

Matt Blumrich  
Dept. of Comp. Science  
Princeton University  
Princeton, NJ 08544

B. W. Boehm  
DARPA/ISTO  
1400 Wilson Blvd.  
Arlington, VA 22209

R. R. Borchert  
L-669  
Lawrence Livermore Nat'l Labs  
PO Box 808  
Livermore, CA 94550

Dan Bowlus  
Mail Code 8123  
Attn: D. Bowlus & G. Letiecq  
Naval Underwater Sys. Cntr.  
Newport, RI 02841-5047

Dr. Donald Brand  
MS - N8930  
Geodynamics

Falcon AFB, CO 80912-5000

Bud Brewster  
410 South Pierce  
Wheaton, IL 60187

Carl N. Brooks  
Brooks Associates  
414 Falls Road  
Chagrin Falls, OH 44022

John Brunet  
245 First St.  
Cambridge, MA 02142

John Bruno  
Cntr for Comp. Sci and Engr  
College of Engineering  
University of California  
Santa Barbara, CA 93106-5110

H. L. Buchanan  
DARPA/DSO  
1400 Wilson Blvd.  
Arlington, VA 22209

D. A. Buell  
Supercomputing Resch. Cntr.  
17100 Science Dr.  
Bowie, MD 20715

B. L. Buzbee  
Scientific Computing Dept.  
NCAR  
PO Box 3000  
Boulder, CO 80307

G. F. Carey  
Dept. of Engineering Mechanics  
TICOM ASE-EM WRW 305  
University of Texas at Austin  
Austin, TX 78712

Art Carlson  
NOSC Code 41  
New London, CT 06320

Bonnie C. Carroll  
Sec. Dir.  
CENDI, Information Int'l  
PO Box 4141  
Oak Ridge, TN 37831

Charles T. Casale  
Aberdeen Group, Inc.  
92 State Street  
Boston, MA 02109

J. M. Cavallini  
US Department of Energy  
OSC, ER-30, GTN  
Washington, DC 20585

John Champine  
Clay Res. Inc., Software Div.  
655F Lone Oak Dr.  
Eagan, MN 55121

Tony Chan  
Department of Computer Science  
The Chinese University of Hong Kong  
Shatin, NT  
Hong Kong

J. Chandra  
Army Research Office  
PO Box 12211  
Resch Triangle Park, NC 27709

Siddhartha Chatterjee  
RIACS  
NASA Ames Research Center  
Mail Stop T045-1  
Moffett Field, CA 94035-1000

Warren Chernock  
Scientific Advisor DP-1  
US Department of Energy  
Forestal Bldg. 4A-045  
Washington, DC 20585

R.C.Y. Chin  
L-321  
Lawrence Livermore Nat'l Lab  
PO Box 808  
Livermore, CA 94550

Mark Christon  
L-122  
Lawrence Livermore Nat'l Lab  
PO Box 808  
Livermore, CA 94550

M. Ciment  
Adv. Sci. Comp/ Div. RM 417  
National Science Foundation  
Washington, DC 20550

Richard Claytor  
US Department of Energy  
Def. Prog. DE-1  
Forestal Bldg. 4A-014  
Washington, DC 20585

Andy Cleary  
Centre for Information Research  
Australia National University  
GPA Box 4  
Canberra, ACT 2601  
Australia

T. Cole  
MS180-500  
Jet Prop. Lab  
4800 Oak Grove Dr.,  
Pasadena, CA 91109

Monte Coleman  
US Army Bal. Resch Lab  
SLCBLR-SE-A (Bldg. 394/216)  
Aberdeen Prov. Grnd., MD 21005-5006

Tom Coleman  
Dept. of Computer Science  
Upson Hall  
Cornell University  
Ithaca, NY 14853

S. Colley  
NCUBE  
19A Davis Drive  
Belmont, CA 94002

J. Coronas  
Ames Laboratory  
236 Wilhelm Hall  
Iowa State University  
Ames, IA 50011-3020

Steve Cosgrove  
E6-220  
Knolls Atomic Power Lab  
PO Box 1072  
Schenectady, NY 12301-1072

C. L. Crothers

IBM Corporation  
472 Wheelers Farms Road  
Milford, CT 06460

J. K. Cullum  
Thomas J. Watson Resch. Center  
PO Box 218  
Yorktown Heights, NY 10598

Leo Dagum  
Computer Sciences Corp.  
NASA Ames Research Center  
Moffett Field, CA 94035

Kenneth I. Daugherty  
Chief Scientist  
HQ DMA (SA), MS -A-16  
8613 Lee Highway  
Fairfax, VA 22031-2138

L. Davis  
Cray Research Inc.  
1168 Industrial Blvd.  
Chippawa Falls, WI 54729

Mr. Frank R. Deis  
Martin Marietta  
Falcon AFB, CO 80912-5000

R. A. DeMillo  
Comp. & Comput. Resch.,  
Rm. 304  
National Science Foundation  
Washington, DC 20550

L. Deng  
Applied Mathematics Dept.  
SUNY at Stony Brook  
Stony Brook, NY 11794-3600

A. Trent DePersia  
Prog. Mgr.  
DARPA/ASTO  
1400 Wilson Blvd.  
Arlington, VA 22209-2308

Sean Dolan  
nCUBE  
919 E. Hillsdale Blvd.  
Foster City, CA 94404

Jack Dongarra  
Department of Computer Science  
University of Tennessee  
Knoxville, TN 37996

L. Dowdy  
Computer Science Department  
Vanderbilt University  
Nashville, TN 37235

Joanne Downey-Burke  
8030 Sangor Dr.  
Colorado Springs, CO 80920

I. S. Duff  
CSS Division  
Harwell Laboratory  
Oxfordshire, OX11 0RA  
United Kingdom

Alan Edelman  
University of California, Berkeley  
Dept. of Mathematics  
Berkeley, CA 94720

S. C. Eisenstat  
Computer Science Dept.

Yale University  
PO Box 2158  
New Haven, CT 06520

H. Elman  
Computer Science Dept.  
University of Maryland  
College Park, MD 20842

M. A. Elmer  
DARPA/RMO  
1400 Wilson Blvd.  
Arlington, VA 22209

J. N. Entzminger  
DARPA/TTO  
1400 Wilson Blvd.  
Arlington, VA 22209

A. M. Erisman  
MS 7L-21  
Boeing Computer Services  
PO Box 24346  
Seattle, WA 98124-0346

R. E. Ewing  
Mathematics Dept.  
University of Wyoming  
PO Box 3036 University Station  
Laramie, WY 82071

El Dabaghi Fadi  
Charge of Research  
Institute Nat'l De Recherche en  
Informatique et en Automatique  
Domaine de Voluceau Rocouencourt  
BP 105  
78153 Le Chesnay Cedex (France)

H. D. Fair  
Institute for Adv. Tech.  
4032-2 W. Braker Lane  
Austin, TX 78759

Kurt D. Fickie  
US Army Ballistic Resch. Lab  
ATTN: SLCBR-SE  
Aberdeen Proving Ground, MD  
21005-5066

Tom Finnegan  
NORAD/USSPACECOM J2XS  
STOP 35  
PETERSON AFB, CO 80914

J. E. Flaherty  
Computer Science Dept.  
Rensselaer Polytech Inst.  
Troy, NY 12181

L. D. Fosdick  
University of Colorado  
Computer Science Department  
Campus Box 430  
Boulder, CO 80309

G. C. Fox  
Northeast Parallel Archit. Cntr.  
111 College Place  
Syracuse, NY 13244

R. F. Freund  
NRAd- Code 423  
San Diego, CA 99152-5000

Sverre Froyen  
Solar Energy Research Inst.  
1617 Cole Blvd.

Golden, CO 80401

David Gale  
Intel Corp  
600 S. Cherry St.  
Denver, CO 80222-1801

D. B. Gannon  
Computer Science Dept.  
Indiana University  
Bloomington, IN 47401

C. W. Gear  
NEC Research Institute  
4 Independence Way  
Princeton, NJ 08548

J. A. George  
Needles Hall  
University of Waterloo  
Waterloo, Ont., Can.  
N2L 3G1

Shomit Ghose  
nCUBE  
919 E. Hillsdale Blvd.  
Foster City, CA 94404

Clarence Giese  
8135 Table Mesa Way  
Colorado Springs, CO 80919

Dr. Horst Gietl  
nCUBE GmbH  
Hanauer Strasse 85  
8000 Munich 50  
Germany

John Gilbert  
Xerox PARC  
3333 Coyote Hill Road  
Palo Alto, CA 94304

Michael E. Giltrud  
DNA  
HQ DNA/SPSD  
6801 Telegraph Rd.  
Alexandria, VA 22310-3398

Alastair M. Glass  
AT&T Bell Labs Rm 1A-164  
600 Mountain Avenue  
Murray Hill, NJ 07974

J. G. Glimm  
Dept. of App Math. & Stat.  
State U. of NY at Stony Brook  
Stony Brook, NY 11794

Dr. Raphael Gluck  
TRW-DSSG, R4/1408  
One Space Park  
Redondo Beach, CA 90278

G. H. Golub  
Computer Science Dept.  
Stanford University  
Stanford, CA 94305

Marcia Grabow  
AT&T Bell Labs 1D-153  
600 Mountain Ave.  
PO Box 636  
Murray Hill, NJ 07974-0636

Nancy Grady  
MS 6240  
Oak Ridge Nat'l Lab

Bos 2008  
Oak Ridge, TN 37831

Anne Greenbaum  
New York University  
Courant Institute  
251 Mercer Street  
New York, NY 10012-1185

Satya Gupta  
Intel SSD  
Bldg. C)6-09, Zone 8  
14924 NW Greenbriar Pkwy  
Beaverton, OR 97006

J. Gustafson  
Computer Science Department  
236 Wilhelm Hall  
Iowa State University  
Ames, IA 50011

R. Grayson Hall  
USDOE/HQ  
1000 Independence Ave, SW  
Washington, DC 20585

Cush Hamlen  
Minnesota Supercomputer Cntr.  
1200 Washington Ave. So.  
Minneapolis, MN 55415

Steve Hammond,  
NCAR  
PO Box 3000  
Boulder, CO 80307

CDR. D. R. Hamon  
Chief, Space Integration Div.  
USSPACECOM/J5SI  
Peterson AFB, CO 80914-5003

Dr. James P. Hardy  
NTBIC/GEODYNAMICS  
MS N 8930  
Falcon AFB, CO 80912-5000

Doug Harless  
NCUBE  
2221 East Lamar Blvd., Suite 360  
Arlington, TX 76006

Mike Heath  
University of Illinois  
4157 Beckman Institute  
405 N. Mathews Ave.  
Urbana, IL 61801

Greg Heileman  
EECE Department  
University of New Mexico  
Albuquerque, NM 87131

Brent Henrich  
Mobile R & D Laboratory  
13777 Midway Rd.  
PO Box 819047  
Dallas, TX 75244-4312

Michael Heroux  
Cray Research Park  
655F Lone Oak Drive  
Eagan, MN 55121

A.J. Hey  
University of Southampton  
Dept. of Electronics and Computer Science  
Mountbatten Bldg., Highfield  
Southampton, SO9 5NH

United Kingdom

W. D. Hillis  
Thinking Machines, Inc.  
245 First St.  
Cambridge, MA 02139

Dan Hitchcock  
US Department of Energy  
SCS, ER-30 GTN  
Washington, DC 20585

LTC Richard Hochbewrg  
SDIO/SDA  
The Pentagon  
Washington, DC 20301-7100

C. J. Holland  
Director  
Math and Information Sciences  
AFOSR/NM, Bolling AFB  
Washington, DC 20332-6448

Dr. Albert C. Holt  
Ofc. of Munitions  
Ofc of Sec. of St.-ODDRE/TWP  
Pentagon, Room 3B1060  
Washington, DC 203301-3100

Mr. Daniel Holtzman  
Vanguard Research Inc.  
10306 Eaton Pl., Suite 450  
Fairfax, VA 22030-2201

David A. Hopkins  
US Army Ballistic Resch. Lab.  
Attention: SLCBR-IB-M  
Aberdeen Prov. Grnd., MD 21005-5066

Graham Horton  
Universitat Erlangen-Nurnberg  
IMMD III  
Martensstrasse 3  
8520 Erlangen  
Germany

Fred Howes  
US Department of Energy  
OSC, ER-30, GTN  
Washington, DC 20585

Chua-Huang Huang  
Assist. Prof. Dept. Comp. & Info Sci  
Ohio State Univ.  
228 Bolz Hall-2036 Neil Ave.  
Columbus, OH 43210-1277

R. E. Huddleston  
L-61  
Lawrence Livermore Nat'l Lab  
PO Box 808  
Livermore, CA 94550

Zdenek Johan  
Thinking Machines Corp.  
245 First Street  
Cambridge, MA 02142-1264

Gary Johnson  
US Department of Energy  
SCS, ER-30 GTN  
Washington, DC 20585

S. Lennart Johnsson  
Thinking Machines Corp.  
245 First Street  
Cambridge, MA 02142-1264

G. S. Jones  
Tech Program Support Cntr.  
DARPA/AVSTO  
1515 Wilson Blvd.  
Arlington, VA 22209

T. H. Jordan  
Dept of Earth, Atmos & Pla. Sci.  
MIT  
Cambridge, MA 02139

M. H. Kalos  
Cornell Theory Center  
514A Eng. and Theory Center  
Hoy Road, Cornell University  
Ithaca, NY 14853

H. G. Kaper  
Math. and Comp. Sci. Division  
Argonne National Laboratory  
Argonne, IL 60439

S. Karin  
Supercomputing Department  
9500 Gilman Drive  
University of CA at San Diego  
La Jolla, CA 92093

Herb Keller  
Applied Math 217-50  
Cal Tech  
Pasadena, CA 91125

M. J. Kelley  
DARPA/DMO  
1400 Wilson Blvd.  
Arlington, VA 22209

K. W. Kennedy  
Computer Science Department  
Rice University  
PO Box 1892  
Houston, TX 77251

Aram K. Kevorkian  
Codje 7304  
Naval Ocean Systems Center  
271 Catalina Blvd.  
San Diego, CA 92152-5000

John E. Killough  
University of Houston  
Dept. of Chem. Engineering  
Houston, TX 77204-4792

D. R. Kincaid  
Cntr. for Num. Analy.,  
RLM 13-150  
University of Texas at Austin  
Austin, TX 78712

T. A. Kitchens  
US Department of Energy  
OSC, ER-30, GTN  
Washington, DC 20585

Thomas Klemas  
394 Briar Lane  
Newark, DE 19711

Dr. Peter L. Kneppell  
NTBIC/GEODYNAMICS  
MS N 8930  
Falcon AFB, CO 80912-5000

Max Koontz  
DOE/OAC/DP 5.1  
Forestal Bldg

1000 Independence Ave.  
Washington, DC 20585

Ann Krause  
HQ AFOTEC/OAN  
Kirtland AFB, NM 87117-7001

V. Kumar  
Computer Science Department  
University of Minnesota  
Minneapolis, MN 55455

J. Lannutti  
MS B-186  
Director, SC. Resch. Institute  
Florida State University  
Tallahassee, FL 32306

P. D. Lax  
New York University-Courant  
251 Mercer St.  
New York, NY 10012

Lawrence A. Lee  
NC Supercomputing Center  
PO Box 12889  
3021 Cornwallis Rd.  
Research Triangle Park, NC 27709

Dr. H.R. Leland  
Calspan Corporation  
PO Box 400,  
Buffalo, NY 14225

David Levine  
Math & Comp. Science  
Argonne National Laboratory  
9700 Cass Avenue South  
Argonne, IL 60439

Peter Littlewood  
Theoret. Phy. Dept.  
AT&T Bell Labs  
Rm 1D-335  
Murray Hill, NJ 07974

Peter Lomdahl  
T-11, MS B262  
Los Alamos Nat'l Lab  
Los Alamos, NM 87545

Louis S. Lome  
SDIO/TNI  
The Pentagon  
Washington, DC 20301-7100

Col. Gordon A. Long  
Deputy Director for Adv. Comp.  
HQ USSPACECOM/JOSDEPS  
Peterson AFB, CO 80914-5003

John Lou  
3258 Caminito Ameca  
La Jolla, CA 92037

Daniel Loyens  
Koninklijke/Shell-Laboratorium  
Postbus 3003  
1002 AA Amsterdam  
The Netherlands

Robert E. Lynch  
Dept. of CS  
Purdue University  
West Lafayette, IN 47907

Kathy MacLeod  
AFEWC/SAT

Kelly AFB  
San Antonio, TX 78243-5000

H. Mair  
Naval Surface Warfare Center  
10901 New Hampshire Ave.  
Silver Springs, MD 20903-5000

Henry Makowitz  
MS - 2211-INEL  
EG&E Idaho Incorporated  
Idaho Falls, ID 83415

David Mandell  
MS F663  
Hydrodynamic App. Grp. X-3  
Los Alamos Nat'l Labs  
Los Alamos, NM 87545

T. A. Manteuffel  
Department of Mathematics  
University of Co. at Denver  
Denver, CO 80202

William S. Mark  
Lockheed - Org. 96-01  
Bldg. 254E  
3251 Hanover Street  
Palo Alto, CA 94303-1191

Kapil Mathur  
Thinking Machines Corporation  
245 First Street  
Cambridge, MA 02142-1214

John May  
Kaman Sciences Corporation  
1500 Garden of the Gods Road  
Colorado Springs, CO 80933

William McColl  
Oxford Univ. Computing Lab  
8-11 Keble Road  
Oxford, OX1 3QD  
United Kingdom

S. F. McCormick  
Computer Mathematics Group  
University of CO at Denver  
1200 Larimer St.  
Denver, CO 80204

J. R. McGraw  
L-316  
Lawrence Livermore Nat'l Lab  
PO Box 808  
Livermore, CA 94550

Jill Mesirov  
Thinking Machines Corporation  
245 First Street  
Cambridge, MA 02142-1214

P. C. Messina  
158-79  
Mathematics & Comp Sci. Dept.  
Caltech  
Pasadena, CA 91125

Prof. Ralph Metcalfe  
Dept. of Mech. Engr.  
University of Houston  
4800 Calhoun Road  
Houston, TX 77204-4792

G. A. Michael  
L-306  
Lawrence Livermore Nat Lab

PO Box 808  
Livermore, CA 94550

Loren K. Miller  
Goodyear Tire & Rubber  
PO Box 3531  
Akron, OH 44309-3531

Robert E. Millstein  
TMC  
245 First Street  
Cambridge, MA 02142

G. Mohnkern  
NOSC - Code 73  
San Diego, CA 92152-5000

C. Moler  
The Mathworks  
24 Prime Park Way  
Natick, MA 01760

Gary Montry  
Southwest Software  
11812 Persimmon, NE  
Albuquerque, NM 87111

N. R. Morse  
C-DO, MS B260  
Comp. & Comm. Division  
Los Alamos National Lab  
Los Alamos, NM 87545

J. R. Moulic  
IBM  
Thomas J. Watson Resch Cntr.  
PO Box 704  
Yorktown Heights, NY 10598

D. B. Nelson  
US Department of Energy  
OSC, ER-30, GTN  
Washington, DC 20585

Jeff Newmeyer  
Org. 81-04, Bldg. 157  
1111 Lockheed Way  
Sunnyvale, CA 94089-3504

D. M. Nosenchuck  
Mech. and Aero. Engr. Dept.  
D302 E Quad  
Princeton University  
Princeton, NJ 08544

C. E. Oliver  
Off. of Lab Comp Bldg. 4500N,  
Oak Ridge Nat'l Laboratory  
PO Box 2008  
Oak Ridge, TN 37831-6259

Dennis L. Orphal  
Calif Resch & Technology Inc.  
5117 Johnson Dr.  
Pleasanton, CA 94588

J. M. Ortega  
Applied Math Department  
University of Virginia  
Charlottesville, VA 22903

John Palmer  
TMC  
245 First St.  
Cambridge, MA 02142

Robert J. Paluck  
Convex Computer Corp.

3000 Waterview Parkway  
PO Box 733851  
Richardson, TX 75083-3851

Anthony C. Parmee  
Counsellor and Attache (Def.)  
British Embassy  
3100 Mass. Ave, NW  
Washington, DC 20008

S. V. Parter  
Department of Mathematics  
Van Vleck Hall  
University of Wisconsin  
Madison, WI 53706

Dr. Nisheeth Patel  
US Army Ballistic Resch. Lab.  
AMXBR-LFD  
Aberdeen Prov. Grnd., MD 21005-5066

A. T. Patera  
Mechanical Engineering Dept.  
77 Massachusetts Ave.  
MIT  
Cambridge, MA 02139

A. Patrino  
Atmos. and Clim. Resch. Div  
Office of Engy Resch, ER-74  
US Department of Energy  
Washington, DC 20545

R. F. Peierls  
Math. Sciences Group, Bldg. 515  
Brookhaven National Lab  
Upton, NY 11973

Donna Perez  
NOSC - MCAC Resource Cntr.  
Code 912  
San Diego, CA 92152-5000

K. Perko  
Supercomputing Solutions, Inc.  
6175 Nancy Ridge Dr.  
San Diego, CA 92121

John Petresky  
Ballistic Research Lab  
SLCBBR-LF-C  
Aberdeen Prov. Grnd., MD 21005-5006

Linda Petzold  
L-316  
Lawrence Livermore Nat'l Lab.  
Livermore, CA 94550

Wayne Pfeiffer  
San Diego SC Center  
PO Box 85608  
San Diego, CA 92186

Frank X. Pfenneberger  
Martin Marietta  
MS-N83104  
National Test Bed  
Falcon AFB, CO 80912-5000

Dr. Leslie Pierre  
SDIO/ENA  
The Pentagon  
Washington, DC 20301-7100

Paul Plassman  
Math and Computer Science Division  
Argonne National Lab  
Argonne, IL 60439

R. J. Plemmons  
Dept. of Math. & Comp Sci.  
Wake Forest University  
PO Box 7311  
Winston Salem, NC 27109

Alex Pothen  
Computer Science Department  
University of Waterloo  
Waterloo, Ontario N2L 3G1  
Canada

John K. Prentice  
Amparo Corporation  
3700 Rio Grande NW, Suite 5  
Albuq., NM 87107-3042

Peter P. F. Radkowski  
PO Box 1121  
Los Alamos, NM 87544

J. Rattner  
Intel Scientific Computers  
15201 NW Greenbriar Pkwy.  
Beaverton, OR 97006

J. P. Retelle  
Org. 94-90  
Lockheed - Bldg. 254G  
3251 Hanover Street  
Palo Alto, CA 94304

C. E. Rhoades  
L-298  
Computational Physics Div.  
PO Box 808  
Lawrence Livermore Nat'l Lab  
Livermore, CA 94550

J. R. Rice  
Computer Science Dept.  
Purdue University  
West Lafayette, IN 47907

John Richardson  
Thinking Machines Corporation  
245 First Street  
Cambridge, MA 02142-1214

Lt. Col. George E. Richie  
Chief, Adv. Tech Plans  
JOSDEPS  
Peterson AFB, CO 80914-5003

John Rollett  
Oxford University Computing Laboratory  
8-11 Keble Road  
Oxford, OX1 3QD  
United Kingdom

R. Z. Roskies  
Physics and Astronomy Dept.  
100 Allen Hall  
University of Pittsburgh  
Pittsburg, PA 15206

Diane Rover  
Michigan State University  
Dept. of Electrical Engineering  
260 Engineering Bldg.  
East Lansing, MI 48824

Y. Saad  
University of Minnesota  
4-192 EE/CSci Bldg.  
200 Union St.  
Minneapolis, MN 55455-0159

P. Sadayappan  
Dept. of Comp. & Info Science  
Ohio State Univ.-228 Bolz Hall  
2036 Neil Ave.  
Columbus, OH 43210-1277

Joel Saltz  
Computer Science Department  
A.V. Williams Building  
University of Maryland  
College Park, MD 20742

A. H. Sameh  
CSR  
305 Talbot Laboratory  
University of Illinois  
104 S. Wright St.  
Urbana, IL 61801

P. E. Saylor  
Dept. of Comp. Science  
222 Digital Computation Lab  
University of Illinois  
Urbana, IL 61801

LCDR Robert J. Schoppe  
Chief, Operations Rqmts  
USSPACECOM/JOSDEPS (Stop 35)  
Peterson AFB, CO 80914

Rob Schreiber  
RIACS  
NASA Ames Research Center  
Mail Stop T045-1  
Moffett Field, CA 94035-1000

M. H. Schultz  
Department of Computer Science  
Yale University  
PO Box 2158  
New Haven, CT 06520

Dave Schwartz  
NOSC, Code 733  
San Diego, CA 92152-5000

Mark Seager  
LLNL, L-80  
PO box 803  
Livermore, CA 94550

A. H. Sherman  
Sci. Computing Assoc. Inc.  
Suite 307, 246 Church Street  
New Haven, CT 06510

Horst Simon  
NAS Systems Division  
NASA Ames Research Center  
Mail Stop T045-1  
Moffett Field, CA 94035

Richard Sincovec  
Oak Ridge National Laboratory  
P.O. Box 2008, Bldg 6012  
Oak Ridge, TN 37831-6367

Vineet Singh  
HP Labs, Bldg. 1U, MS 14  
1501 Page Mill Road  
Palo Alto, CA 94304

Ashok Singhal  
CFD Resch. Center  
3325 Triana Blvd.  
Huntsville, AL 35805

Anthony Skjellum  
Lawrence Livermore National Laboratory  
7000 East Ave., Mail Code L-316  
Livermore, CA 94550

L. Smarr  
Director, Supercomputer Apps.  
152 Supercomputer Applications  
Bldg. 605 E. Springfield  
Champaign, IL 61801

William R. Somsay  
Ballistic Research Laboratory  
SLCBBR-SE-A, Bldg. 394  
Aberdeen Proving Ground, MD 21005

D. C. Sorenson  
Department of Math Sciences  
Rice University  
PO Box 1892  
Houston, TX 77251

S. Squires  
DARPA/ISTO  
1400 Wilson Blvd.  
Arlington, VA 11109

N. Srinivasan  
AMOCO Corp  
PO 87703  
Chicago, IL 60680-0703

Thomas Stegmann  
Digital Equipment Corporation  
8085 S. Chester Street  
Englewood, CO 80112

D. E. Stein  
AT&T  
100 South Jefferson Rd.  
Whippany, NJ 07981

M. Steuerwalt  
Division of Math Sciences  
National Science Foundation  
Washington, DC 20550

Mike Stevens  
nCUBE  
919 E. Hillsdale Blvd.  
Foster City, CA 94404

G. W. Stewart  
Computer Science Department  
University of Maryland  
College Park, MD 20742

O. Storassli  
MS-244  
NASA Langley Research Cntr.  
Hampton, VA 23665

C. Stuart  
DARPA/TTO  
1400 Wilson Blvd.  
Arlington, VA 22209

LTC James Sweeder  
SDIO/SDA  
The Pentagon  
Washington, DC 20301-7100

R. A. Tapia  
Mathematical Sci. Dept  
Rice University  
PO Box 1892  
Houston, TX 77251

Gligor A. Tashkovich  
PO Box 296  
Pound Ridge, NY 10576-0296

H. Teuteberg  
Cray Research, Suite 830  
6565 Americas Pkway, NE  
Albuquerque, NM 87110

A. Thaler  
Division of Math Sciences  
National Science Foundation  
Washington, DC 20550

Allan Torres  
125 Lincoln Ave., Suite 400  
Santa Fe, NM 87501

Harold Trease  
Los Alamos National Lab  
PO Box 1666, MS F663  
Los Alamos, NM 87545

Randy Truman  
Mechanical Engineering Dept.  
University of NM  
Albuquerque, NM 87131

Ray Tuminaro  
CERFACS  
42 Ave Gustave Coriolis  
31057 Toulouse Cedex  
France

Mark Urrutia  
Intel Corp., CO1-03  
5200 NE Elam Young Pkwy.  
Hillsboro, ORE 97124-6497

Mike Uttormark  
UW-Madison  
1500 Johnson Dr.  
Madison, WI 53706

R. VanDeGeijn  
Computer Science Department  
University of Texas  
Austin, TX 78712

George Vandergrift  
Dist. Mgr.  
Convex Computer Corp.  
3916 Juan Tabo, NE, Suite 38  
Albuquerque, NM 87111

H. VanDerVorst  
Delft University of Technology  
Faculty of Mathematics  
POB 356  
2600 AJ Delft  
The Netherlands

C. VanLoan  
Department of Computer Science  
Cornell University, Rm. 5146  
Ithaca, NY 14853

John VanRosendale  
ICASE, NASA Langley Research Center  
MS 132C  
Hampton, VA 23665

Steve Vavasis  
Dept. of Computer Science  
Upson Hall  
Cornell University  
Ithaca, NY 14853

R. G. Voigt  
MS 132-C  
NASA Langley Resch Cntr, ICASE  
Hampton, VA 36665

Phuong Vu  
Cray Research, Inc.  
19607 Franz Road  
Houston, TX 77084

David Walker  
Bldg 6012  
Oak Ridge National Lab  
PO Box 2008  
Oak Ridge, TN 37831

Steven J. Wallach  
Convex Computer Corp.  
3000 Waterview Parkway  
PO Box 833851  
Richardson, TX 75083-3851

R. C. Ward  
Bld. 9207-A  
Mathematical Sciences  
Oak Ridge National Lab  
PO Box 4141  
Oak Ridge, TN 37831-8083

Thomas A. Weber  
National Science Found.  
1800 G. Street, NW  
Washington, DC 20550

G. W. Weigand  
DARPA/CSTO  
3701 N. Fairfax Ave.  
Arlington, VA 22203-1714

M. F. Wheeler  
Math Sciences Dept  
Rice University  
PO Box 1892  
Houston, TX 77251

A. B. White  
MS-265  
Los Alamos National Lab  
PO Box 1663  
Los Alamos, NM 87544

B. Wilcox  
DARPA/DSO  
1400 Wilson Blvd.  
Arlington, VA 22209

Roy Williams  
California Institute of Technology  
206-49  
Pasadena, CA 91104

C. W. Wilson  
MS MI02-3/B11  
Digital Equipment Corp.  
146 Main Street  
Maynard, MA 00175

K. G. Wilson  
Physics Dept.  
Ohio State University  
Columbus, OH 43210

Leonard T. Wilson  
NSWC  
Code G22  
Dahlgren, VA 22448

Peter Wolochow

Intel Corp., CO1-03  
5200 NE Elam Young Pkwy.  
Hillsboro, OR 97124-6497

P. R. Woodward  
University of Minnesota  
Department of Astronomy  
116 Church Street, SE  
Minneapolis, MN 55455

M. Wunderlich  
Math. Sciences Program  
National Security Agency  
Ft. George G. Mead, MD 20755

Hishashi Yasumori  
KTEC-Kawasaki Steel  
Techno-research Corporation  
Hibiya Kokusai Bldg. 2-3  
Uchisaiwaicho 2-chrome  
Chiyoda-ku, Tokyo 100

David Young  
Center for Numerical Analysis  
RLM 13.150  
The University of Texas  
Austin, TX 78713-8510

Robert Young  
Alcoa Laboratories  
Alcoa Center, PA 15069  
Attn: R. Young & J. McMichael

William Zierke  
Applied Research Lab-Penn State.  
PO Box 30  
State College, PA 16804

#### INTERNAL DISTRIBUTION:

Paul Fleury	1000
Ed Barsis	1400
Sudip Dosanjh	1402
Bill Camp	1421
Doug Cline	1421
David Gardner	1421
Grant Heffelfinger	1421
Scott Hutchinson	1421
Martin Lewitt	1421
Steve Plimpton	1421
Mark Sears	1421
John Shadid	1421
Julie Swisshelm	1421
Dick Allen	1422
Bruce Hendrickson (25)	1422
David Womble	1422
Ernie Brickell	1423
Kevin McCurley	1423
Robert Benner	1424
Carl Diegert	1424
Art Hale	1424
Rob Leland (25)	1424
Courtenay Vaughan	1424
Steve Attaway	1425
Johnny Biffle	1425
Mark Blanford	1425
Jim Schutt	1425
Michael McGlaun	1431
Allen Robinson	1431
Paul Yarrington	1432
David Martinez	1434
Dona Crawford	1900
William Mason	1952
Technical Library (5)	7141
Technical Publications	7151
Document Processing for	
DOE/OSTI (10)	7613-2
Central Technical File	8523-2
Charles Tong	8117